


Hack Session

From Language to Robotics: Practical Lessons Bridging LLMs, RL, and AI

Speaker

Logesh Kumar Umapathi


Machine Learning Consultant | Blackbox.ai



Hi 🖐️

I am Logesh Kumar Umapathi

Machine Learning Consultant @ Blackbox.ai



Twitter: @logesh_umapathi


Linkedin: www.linkedin.com/in/logeshkumaru/

Website: logeshumapathi.com

Research : <https://scholar.google.com/citations?user=eWwEqToAAAAJ&hl=en&authuser=1>


Robotics : One of my nine lives


Graham Weaver: <https://www.grahamweaver.com/blog/stanford-graduate-business-school-last-lecture-2024>



Agenda


1. A Glimpse
2. Meet the Hardware - So100
3. Fundamentals of Imitation Learning
4. Popular Policies:
 - a. ACT
 - b. SmolVLA
 - c. Pi0
5. How to record a dataset
6. How not to train the arm
7. Demo





Meet the Hardware


Spec	Value
Degrees of freedom	6 (+ 1 gripper)
Motors	6 serial bus servos
Build	Mostly 3-D printed PLA
Cost	~250 USD per Arm (India: INR 40000)
Strength	200–300 g payload (w/ 12 V 30 kg cm servos)
Reach	~40 cm, $\pm 180^\circ$ sweep



Academic Labs : Cost of building


Bill of Materials

Part	Quantity	Link	Price (per unit)
Robots			
ViperX 300 Robot Arm 6DOF	2	https://www.trossenrobotics.com/viperx-300-robot-arm-6dof.aspx	\$5,695.95
WidowX 250 Robot Arm 6DOF	2	https://www.trossenrobotics.com/widowx-250-robot-arm-6dof.aspx	\$3,295.95
Robot Cage			
Standing Desk (48x30", Black)	1	https://a.co/d/4JuSVhu	\$199.99
Extrusion Bars Black, 10PCS 1220mm(48")	1	https://a.co/d/8ywG7Eu	\$99.99
Extrusion Bars Black, 10PCS 1000mm(39.4")	1	https://a.co/d/8ywG7Eu	\$83.99
Extrusion Bars Black, 4PCS 150mm(5.9")	2	https://a.co/d/8ywG7Eu	\$11.99
Corner Bracket 20Set M5x8	3	https://a.co/d/9BCcU5S	\$28.99
Corner bracket Black L-4 with Screw	4	https://a.co/d/ihrw6YW	\$15.99
Rope Crimping Tool (15 inch for 3/64" to 1/8")	1	https://a.co/d/1t1EszZ	\$35.97



BOM


Item	Qty	Cost (INR)
12v ST3215 Servo	12	24000
Motor Control Board	2	1048
Top USB camera	1	1198
Wrist USB camera	1	3000
Table Clamp 4pcs	1	1101
Power Supply 12v	2	1390
3D Printing Parts		7600
Total		~40,000 INR



BOM for rest of the world and Printing guide: <https://github.com/TheRobotStudio/SO-ARM100/blob/main/SO100.md>

Imitation Learning 101

1. Supervised learning on $\langle \text{image, state} \rangle \rightarrow \langle \text{next action} \rangle$ pairs
2. Advantages: simple, data-efficient, safe (no exploration)
3. Limitations: covariate shift & unseen states \rightarrow compounding errors



What is a policy?


- A policy is a function / model (π) that maps the current state (S) of the robot to an action

$$\pi: S \rightarrow A$$

S - the state is usually the position of the robot, the cameras and sensors feed, and the text instructions.

A - the actions the 6-DOF (degrees of freedom)

cartesian position (x, y, z, rx, ry, rz), the angles of the joints...



Action Chunking Transformer

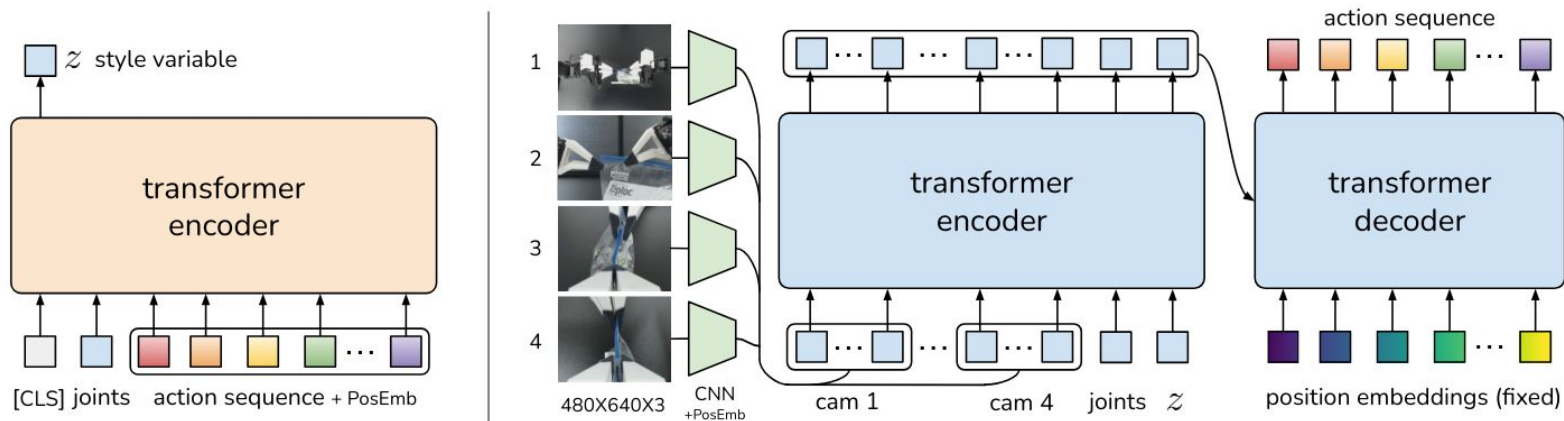




Fig. 4: *Architecture of Action Chunking with Transformers (ACT)*. We train ACT as a Conditional VAE (CVAE), which has an encoder and a decoder. *Left*: The encoder of the CVAE compresses action sequence and joint observation into z , the style variable. The encoder is discarded at test time. *Right*: The decoder or policy of ACT synthesizes images from multiple viewpoints, joint positions, and z with a transformer encoder, and predicts a sequence of actions with a transformer decoder. z is simply set to the mean of the prior (i.e. zero) at test time.




Action Chunking Transformer

Action Chunking




Action Chunking + Temporal Ensemble




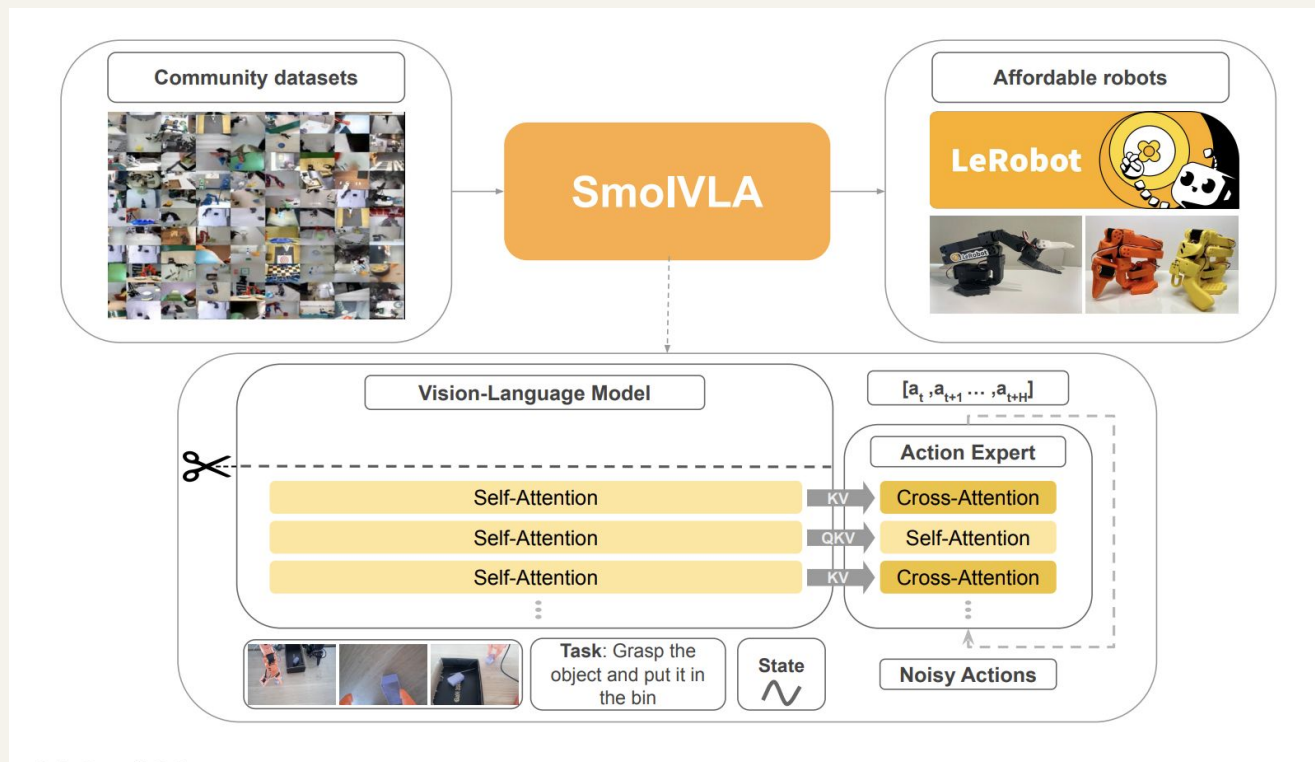
Train

```
python src/lerobot/scripts/train.py \  
    --dataset.repo_id=repo_id/dataset_name \  
\  
    --policy.type=act \  
    --output_dir=outputs/train/output_dir \  
    --job_name=act_sol100_test \  
    --policy.device=mps \  
    --wandb.enable=true \  
    --policy.push_to_hub=False
```





SmoIVLA



SmolVLA

Policy	VLA pt.	Success Rate (%) – Real World			
		Pick-Place	Stacking	Sorting	Avg.
Single-task Training					
SmolVLA (0.45B)	No	55	45	20	40
Multi-task Training					
SmolVLA (0.45B)	No	80	40	35	51.7
SmolVLA (0.45B)	Yes	75	90	70	78.3


community data helps!

Policy	Success Rate (%) – Real World			
	Pick-Place	Stacking	Sorting	Avg.
Single-task Training				
ACT	70	50	25	48.3
Multi-task Training				
π_0 (3.5B)	100	40	45	61.7
SmolVLA (0.45B)	75	90	70	78.3

Table 3 | Real-world benchmarks (SO100). Success rate (%) across three tasks using policies trained in multi-task and single-task settings.


Policy	Success Rate (%) – Real World	
	In Distribution	Out of Distribution
Single-task Training		
ACT	70	40
SmolVLA (0.45B)	90	50

Table 4 | Real-world benchmark (SO101). Success rate (%) for the Pick-Place-Lego task using policies trained in single-task setting.




SmolVLA



- Flow Matching Transformer (~100M parameters) : direct, non-autoregressive prediction of continuous actions
- Visual Token Reduction: 512×512 image is compressed into just 64 tokens, instead of 1024
-
- Faster Inference via Layer Skipping: action expert only attends to VLM features up to half the total layers



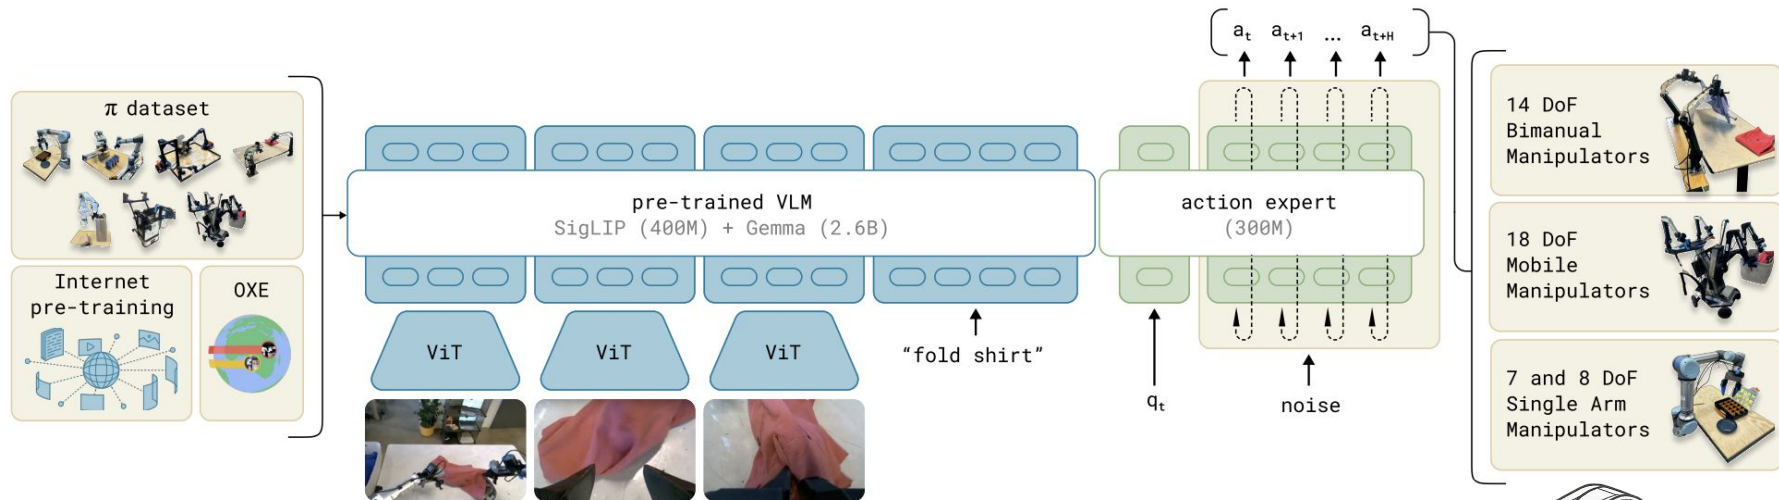
Train

```
python src/lerobot/scripts/train.py \  
    --policy.path=lerobot/smolvla_base \  
    --dataset.repo_id=infinitylogesh/grab_toy_car \  
    --batch_size=128 \  
    --steps=20000 \  
    --save_freq=10000 \  
    --output_dir=outputs/train/grab_toy_car_2_cam_smolvla \  
    --job_name=grab_toy_car_2_cam_smolvla \  
    --policy.device=cuda \  
    --policy.repo_id=infinitylogesh \  
    --wandb.enable=true \  
    --policy.push_to_hub=true
```






Pi0




$$\mathbf{A}_t = [a_t, a_{t+1}, \dots, a_{t+H-1}]$$

$$\mathbf{o}_t = [I_t^1, \dots, I_t^n, \ell_t, q_t]$$




How to record a dataset



Assemble


Follow this guide:

<https://huggingface.co/docs/lerobot/en/so100>




Calibration

```
python -m lerobot.calibrate \  
    --teleop.type=so100_leader \  
    --teleop.port=/dev/tty.usbmodem59591125071 \  
    --teleop.id=my_leader
```



Teleoperation

```
python -m lerobot.teleoperate \  
  --robot.type=so100_follower \  
  --robot.port=/dev/tty.usbmodem58FA0958621 \  
  --robot.id=my_follower \  
  --robot.cameras="{ wrist: {type: opencv, index_or_path: 0, width: 1280, height: 720, fps: 30}, top: \  
    {type: opencv, index_or_path: 1, width: 1280, height: 720, fps: 30}}" \  
  --teleop.type=so100_leader \  
  --teleop.port=/dev/tty.usbmodem59591125071 \  
  --teleop.id=my_leader \  
  --display_data=true
```




Controlling the Arm 🖐️ → 🤖

1. Leader-Follower Teleoperation

- Two identical arms
- Encoders on leader stream joint angles → follower mirrors
- Intuitive, latency-free, but doubles hardware cost


2. Game-pad Teleoperation (PS4 Controller)

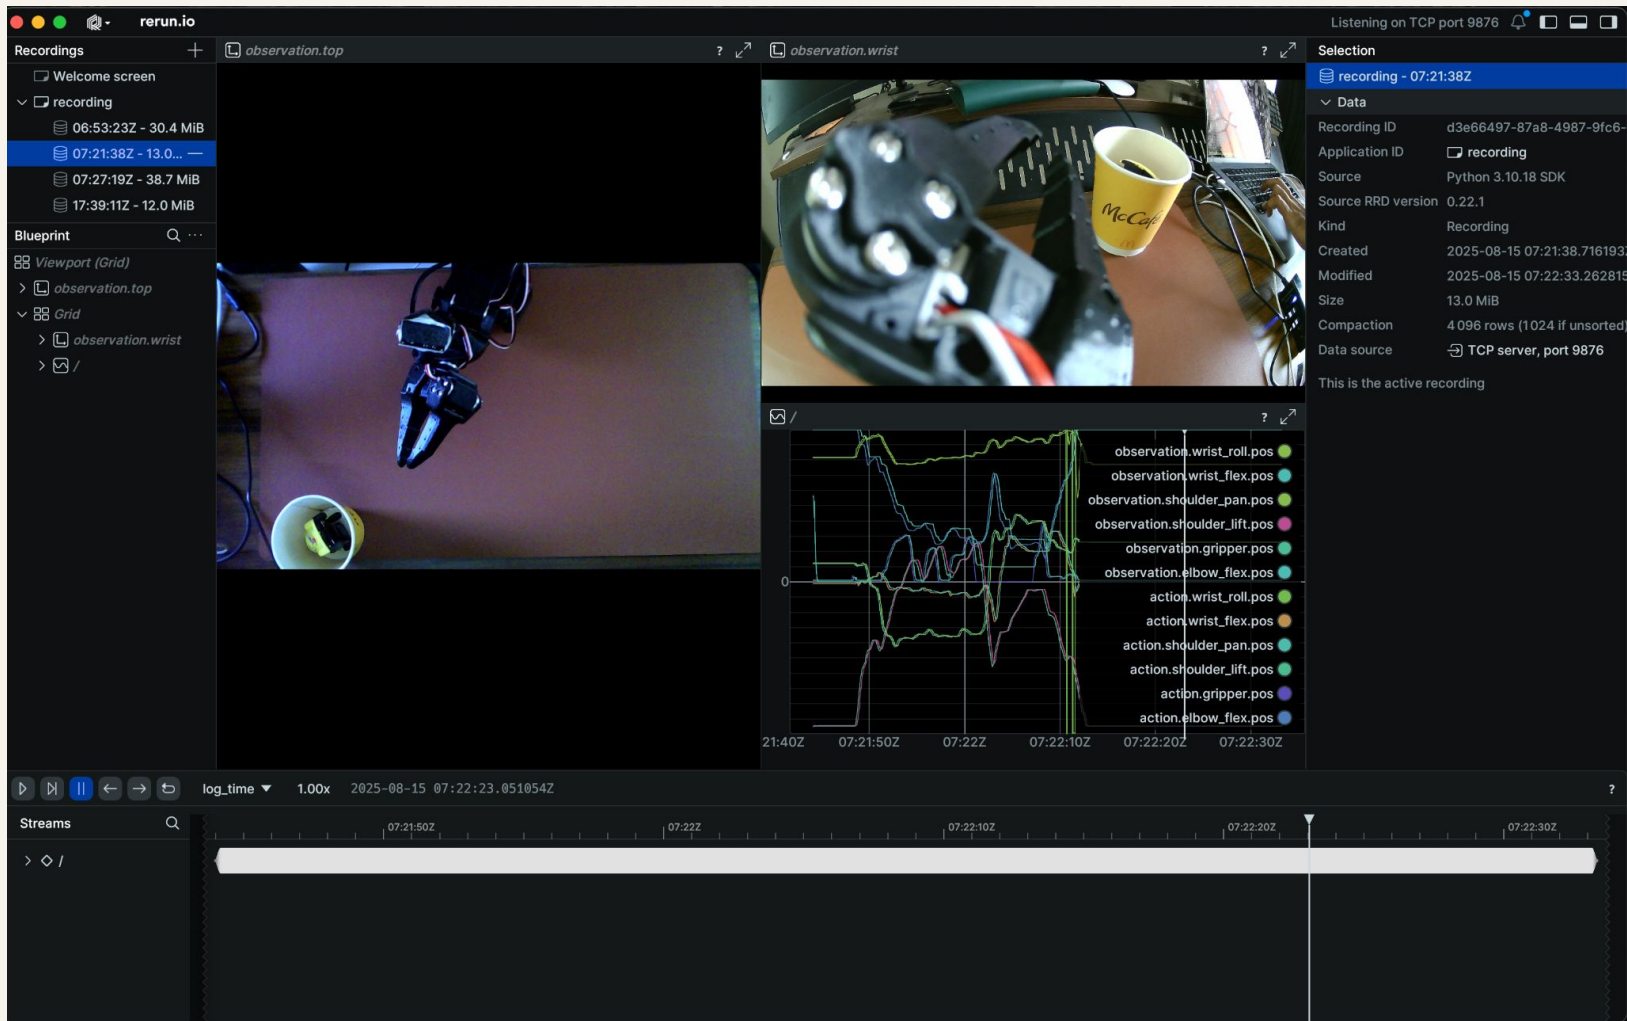
- Map sticks/buttons to joints & gripper
 - Haptic & LED feedback for joint limits
 - Preset poses & accelerometer control
- > Design rule: Match interfaces to your data-collection needs.





Video source : <https://www.youtube.com/shorts/zNo300IixPs>





Action Chunking Transformer

```
[4] from datasets import load_dataset
import pandas as pd

# Load the 'lerobot/pusht' dataset
dataset = load_dataset("lerobot/svla_so101_pickplace", split="train")

# Convert to pandas DataFrame for easier exploration
df = pd.DataFrame(dataset)

# Display the first 5 rows
display(df.head())
```


	action	observation.state	timestamp	frame_index	episode_index	index	task_index
0	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.000000	0	0	0	0
1	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.033333	1	0	1	0
2	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.066667	2	0	2	0
3	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.100000	3	0	3	0
4	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.133333	4	0	4	0

```
df["observation.state"][0]
```

```
[1.9560878276824951,  
-98.74372100830078,  
98.92424774169922,  
74.81983947753906,  
-51.4529914855957,  
1.409399331512451]
```

```
df["action"][0]
```

```
[1.9117647409439087,  
-99.41053009033203,  
99.54483032226562,  
74.84062957763672,  
-48.52258682259766,  
1.6228748559951762]
```




Readings from an episode

```
# first episode
display(df[df.episode_index==0])
```

	action	observation.state	timestamp	frame_index	episode_index	index	task_index
0	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.000000	0	0	0	0
1	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.033333	1	0	1	0
2	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.066667	2	0	2	0
3	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.100000	3	0	3	0
4	[1.9117647409439087, -99.41053009033203, 99.54...	[1.9560878276824951, -98.74372100830078, 98.92...	0.133333	4	0	4	0
...
298	[4.485294342041016, -99.41053009033203, 99.453...	[4.5109782218933105, -98.82746887207031, 98.83...	9.933333	298	0	298	0
299	[4.485294342041016, -99.41053009033203, 99.453...	[4.5109782218933105, -98.82746887207031, 98.83...	9.966666	299	0	299	0
300	[4.485294342041016, -99.41053009033203, 99.453...	[4.5109782218933105, -98.82746887207031, 98.83...	10.000000	300	0	300	0
301	[4.485294342041016, -99.41053009033203, 99.453...	[4.5109782218933105, -98.82746887207031, 98.83...	10.033334	301	0	301	0
302	[4.485294342041016, -99.41053009033203, 99.453...	[4.5109782218933105, -98.82746887207031, 98.83...	10.066667	302	0	302	0

303 rows x 7 columns




Total episode duration: 10 secs

Reading frequency : 30 times per sec (30hz)



Total number of readings : 302

Practical Teleop & Dataset Tips


1. Don't show the leader arm in any training camera view. Otherwise the policy learns to imitate the controller, not the task.
2. Use contrast backgrounds & anti-slip mats to reduce reflection & drift.
3. Diversify positions—cover edges & extremes of the workspace.
4. 50–60 demos \approx enough for a simple pick-and-place skill.
5. Use a wrist camera and train on a stable surface.
6. Good to have a constant and stable source of light.




How not to record a dataset




Demo | Recording a dataset




Demo | Training and Inference - trained policy



Demo | Calibrating the Arms



What's the big deal ?



Thank you



Slides & code:



Feedback & Questions

Twitter: @logesh_umapathi

Linkedin: www.linkedin.com/in/logeshkumaru/

<https://128.pl/1NLbm>